

PCA Revealed

Part 3: Preliminary Concepts

Gaston Sanchez

August 2014

Content licensed under [CC BY-NC-SA 4.0](#)

Readme

License:

Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
<http://creativecommons.org/licenses/by-nc-sa/4.0/>

You are free to:

- Share** — copy and redistribute the material
- Adapt** — rebuild and transform the material

Under the following conditions:

- Attribution** — You must give appropriate credit, provide a link to the license, and indicate if changes were made.
- NonCommercial** — You may not use this work for commercial purposes.
- Share Alike** — If you remix, transform, or build upon this work, you must distribute your contributions under the same license to this one.

Introduction

Introduction

About

Before starting our discussion about **Principal Components Analysis** (PCA), we need to review some conceptual elements of extreme importance.

Main Idea

The idea is to introduce some material that surrounds PCA, but at the same time, this material can be presented independently of PCA.

About the Data

About PCA

Briefly:

Principal Components Analysis (PCA) is a multivariate method that allows us to study and explore a data set of quantitative variables measured on a set of objects.

Data Structure

Analyzed Data

PCA is applied to study a data set under a rectangular format:
table with rows and columns.

Conventionally

- ▶ Rows represent objects (i.e. observations, individuals, samples).
- ▶ Columns represent variables (i.e. features, characteristics, descriptors).

Example: Cereals Data Set

cereals

##	Cups	Calories	Carbs	Fat	Fiber	Potassium	Protein	Sodium	Sugars
## CapnCrunch	0.75	120	12.0	2	0.0	35	1	220	12
## CocoaPuffs	1.00	110	12.0	1	0.0	55	1	180	13
## Trix	1.00	110	13.0	1	0.0	25	1	140	12
## AppleJacks	1.00	110	11.0	0	1.0	30	2	125	14
## CornChex	1.00	110	22.0	0	0.0	25	2	280	3
## CornFlakes	1.00	100	21.0	0	1.0	35	2	290	2
## Nut&Honey	0.67	120	15.0	1	0.0	40	2	190	9
## Smacks	0.75	110	9.0	1	1.0	40	2	70	15
## MultiGrain	1.00	100	15.0	1	2.0	90	2	220	6
## CracklinOat	0.50	110	10.0	3	4.0	160	3	140	7
## GrapeNuts	0.25	110	17.0	0	3.0	90	3	179	3
## HoneyNutCheerios	0.75	110	11.5	1	1.5	90	3	250	10
## NutriGrain	0.67	140	21.0	2	3.0	130	3	220	7
## Product19	1.00	100	20.0	0	1.0	45	3	320	3
## TotalRaisinBran	1.00	140	15.0	1	4.0	230	3	190	14
## WheatChex	0.67	100	17.0	1	3.0	115	3	230	3
## Oatmeal	0.50	130	13.5	2	1.5	120	3	170	10
## Life	0.67	100	12.0	2	2.0	95	4	150	6
## Maypo	1.00	100	16.0	1	0.0	95	4	0	3
## QuakerOats	0.50	100	14.0	1	2.0	110	4	135	6
## Muesli	1.00	150	16.0	3	3.0	170	4	150	11
## Cheerios	1.25	110	17.0	2	2.0	105	6	290	1
## SpecialK	1.00	110	16.0	0	1.0	55	6	230	3

Cereals Data

Download the data in R using the package RCurl as follows:

```
# if you don't have RCurl installed
install.packages("RCurl")

# load package RCurl
library(RCurl)

# google docs spreadsheets url
google_docs = "https://docs.google.com/spreadsheet/"

# public key of data 'cereals'
cereals_key = "pub?key=0AjoVnZ9iB261dEljWDZiWS1reVh1UFNKdE5EcVRvSkE&output=csv"

# download URL of data file
cereals_csv = getURL(paste(google_docs, cereals_key, sep = ""))

# import data in R (through a text connection)
cereals = read.csv(textConnection(cereals_csv), row.names = 1, header = TRUE)
```

Cereals Data

Structure of data cereals

```
# check structure
str(cereals, vec.len = 1)

## 'data.frame': 23 obs. of  9 variables:
##  $ Cups      : num  0.75 1 ...
##  $ Calories  : int  120 110 ...
##  $ Carbs     : num  12 12 ...
##  $ Fat       : int  2 1 ...
##  $ Fiber     : num  0 0 ...
##  $ Potassium: int  35 55 ...
##  $ Protein   : int  1 1 ...
##  $ Sodium    : int  220 180 ...
##  $ Sugars    : int  12 13 ...
```

Example: Cereals Data Set

Variables in Cereals

$X_1 = \text{Cups}$

$X_2 = \text{Calories}$

\vdots

$X_8 = \text{Sodium}$

$X_9 = \text{Sugars}$

Objects in Cereals

$obj_1 = \text{CapnCrunch}$

$obj_2 = \text{CocoaPuffs}$

\vdots

$obj_{22} = \text{Cheerios}$

$obj_{23} = \text{SpecialK}$

Matrix Structure

Data

The analyzed data can be expressed in matrix format \mathbf{X} :

$$\mathbf{X}_{n,p} = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

- ▶ n objects in the rows
- ▶ p quantitative variables in the columns

Variables Considerations

Variables

We will denote the p variables in \mathbf{X} by X_1, X_2, \dots, X_p

Mean-centered

We will assume that the data is centered

$$\bar{X}_j = \sum_{i=1}^n x_{ij} = 0$$

(i.e. *centered*: variables with mean = 0)

Mean-centered Data

Centered Data

For convenient reasons—to make computations easier and simplify notation—it is typically assumed that the variables are **mean-centered** which means that variables have $\text{mean} = 0$.

We often refer to this assumption by simply saying that **the data is centered**.

Centering Variables in Cereals Data

Centered data cereals using function `scale()`

```
# mean-centering variables
X = scale(cereals, center = TRUE, scale = FALSE)

# check mean values of variables
colMeans(X)
```

##	Cups	Calories	Carbs	Fat	Fiber	Potassium	Protein
##	2.896e-17	2.471e-15	6.951e-16	8.689e-17	-4.827e-17	3.707e-15	-5.792e-17
##	Sodium	Sugars					
##	-2.471e-15	3.475e-16					

More about the variables

Variables assumption

When performing PCA, there's an implicit assumption that no single feature or variable is more important than any other.

Interdependent Variables

This implies that we make no distinctions of variables as being predictors or predictands. Some authors say that all variables are **interdependent**.

Looking at Rows and Columns

Data Concerns


Two sides of the same coin

When the analyzed data can be expressed as a matrix with objects in rows, and variables in columns, we commonly care for two issues:

- ▶ Study the **resemblance between objects**
- ▶ Study the **relationships among variables**

Data Perspectives

looking at a data matrix from two perspectives



Data Matrix Products

Matrix Products

There are **two fundamental matrix products** that play a crucial role when the data is in a matrix X with objects in rows, and variables in columns:

- ▶ $X'X$ association matrix for the variables
- ▶ XX' association matrix for the objects

(keep in mind we are assuming centered data)

Association Matrix of Variables

Association between Variables

Having data in a matrix X with n rows and p variables, the association matrix for variables $X'X$ is a matrix of size $p \times p$

Cross-Products of Variables

This matrix contains the cross-products between all pairs of variables. Moreover, this matrix represents the relationships between variables. It is the matrix containing the information of the variances and covariances.

Cross-Products among variables

```
# association matrix of variables
```

```
assoc_variables = t(X) %*% X
```

```
# dimensions
```

```
dim(assoc_variables)
```

```
## [1] 9 9
```

```
# first 5 rows and 5 columns
```

```
assoc_variables[1:5,1:5]
```

```
##           Cups Calories   Carbs      Fat   Fiber
## Cups      1.2779   -1.713    4.402  -0.8791  -2.395
## Calories -1.7130  4486.957   56.957  140.8696  140.435
## Carbs     4.4020   56.957  284.457  -24.6304   1.935
## Fat      -0.8791  140.870 -24.630   18.6087  10.804
## Fiber    -2.3946  140.435   1.935   10.8043  37.152
```

Association Matrix of Objects

Associations between Objects

Having data in a matrix X with n rows and p variables, the association matrix for variables XX' is a matrix of size $n \times n$

Cross-Products of Objects

This matrix contains the cross-products between all pairs of objects. Moreover, this matrix represents the relationships between objects.

Cross-products between objects

```
# association matrix of objects
```

```
assoc_objects = X %*% t(X)
```

```
# dimensions
```

```
dim(assoc_objects)
```

```
## [1] 23 23
```

```
# first 5 rows and 5 columns
```

```
assoc_objects[1:5, 1:5]
```

```
##           CapnCrunch CocoaPuffs   Trix AppleJacks CornChex
## CapnCrunch    3619.3    1325.5 1655.4     958.9    5791.0
## CocoaPuffs    1325.5    1133.8 2462.7    2469.2     990.3
## Trix          1655.4    2462.7 6293.5    6746.1    -760.9
## AppleJacks     958.9    2469.2 6746.1    7459.6   -2442.3
## CornChex      5791.0     990.3 -760.9   -2442.3   11948.7
```


Variation in Data

Total Variation in Data

Variation and Inertia

A major concept is that of **total variation** in data, also known as **total inertia**

About Inertia

Inertia reflects the total amount of dispersion (variability) in the data. Two ways to compute inertia using association matrices:

- ▶ $\text{trace}(X'X)$
- ▶ $\text{trace}(XX')$

Inertia

The inertia is a numeric value that indicates the amount of dispersion in a dataset. Think of it as a generalization of the concept of variance.

By itself the inertia doesn't tell us much, but it's important to know that it can be calculated from both association matrices (objects and variables)

```
# inertia  
sum(diag(assoc_variables))
```

```
## [1] 188340
```

```
# inertia  
sum(diag(assoc_objects))
```

```
## [1] 188340
```

Matrix Decompositions

Decompositions

Decompositions

Matrix decompositions, also known as matrix factorizations, are another topic that plays a relevant role. They allow us to express a matrix \mathbf{Y} as the product of a number of simpler matrices —usually 2 or 3—:

$$\mathbf{Y} = \mathbf{AB}$$

or

$$\mathbf{Y} = \mathbf{ABC}$$

Data Matrix Products

What for?

A matrix decomposition is just a means of expressing a matrix as a product of two or more simpler matrices.

EVD and SVD

There are many types of matrix decompositions but for our PCA understanding purposes, we are interested in two decompositions:

- ▶ Eigen-Value Decomposition (EVD)
- ▶ Singular Value Decomposition (SVD)

EVD

Eigenvalue Decomposition

This decomposition applies to **symmetric** matrices such as the cross-product association matrices $X'X$ and XX'

Symmetric Matrices

The attractive thing about EVD is that when applied to symmetric matrices the results have a “simple” nice structure.

Eigen-Value Decomposition

EVD

An $n \times n$ **symmetric matrix** \mathbf{Y} can be decomposed as:

$$\mathbf{Y} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}'$$

where

- ▶ \mathbf{U} is a $n \times p$ column **orthonormal** matrix containing the eigen-vectors of \mathbf{Y}
- ▶ $\mathbf{\Lambda}$ is a $p \times p$ **diagonal** matrix containing the eigen-values of \mathbf{Y}

EVD in R

`eigen()` function

R provides the function `eigen()` to perform a eigen-value decomposition of a given matrix

`eigen()` output

A list with the following components

- ▶ `values` a vector containing the eigen-values
- ▶ `vectors` a matrix whose columns contain the eigen-vectors

EVD example in R

```
# create symmetric matrix Y
set.seed(22)
M = matrix(rnorm(20), 5, 4)
Y = t(M) %*% M

# eigen-value decomposition
EVD = eigen(Y)

# elements returned by svd()
names(EVD)

## [1] "values" "vectors"

# vector of singular values
(lambda = EVD$values)

## [1] 15.615  4.090  2.175  0.187
```

```
# matrix of eigen-vectors
(U = EVD$vectors)

##           [,1]      [,2]      [,3]      [,4]
## [1,]  0.5708  0.7407 -0.33863  0.1043
## [2,] -0.2742  0.5295  0.76797  0.2338
## [3,]  0.2772 -0.3206  0.04462  0.9046
## [4,]  0.7226 -0.2612  0.54181 -0.3408

# U orthonormal (U'U = I)
t(U) %*% U

##           [,1]      [,2]      [,3]      [,4]
## [1,]  1.000e+00 -3.331e-16 -1.110e-16 -2.776e-17
## [2,] -3.331e-16  1.000e+00  3.608e-16  1.665e-16
## [3,] -1.110e-16  3.608e-16  1.000e+00  5.274e-16
## [4,] -2.776e-17  1.665e-16  5.274e-16  1.000e+00
```

EVD example in R (con't)

```
# Y equals U L U'
U %*% diag(lambda) %*% t(U)

##          [,1]    [,2]    [,3]    [,4]
## [1,]  7.584 -1.401  1.485  5.244
## [2,] -1.401  3.614 -1.767 -2.769
## [3,]  1.485 -1.767  1.778  3.466
## [4,]  5.244 -2.769  3.466  9.092

# compare to Y
Y

##          [,1]    [,2]    [,3]    [,4]
## [1,]  7.584 -1.401  1.485  5.244
## [2,] -1.401  3.614 -1.767 -2.769
## [3,]  1.485 -1.767  1.778  3.466
## [4,]  5.244 -2.769  3.466  9.092
```

SVD

Singular Value Decomposition

This decomposition applies for any rectangular matrix, meaning that we can apply SVD to any data table X .

Principal Axis

The real importance behind SVD comes from the SVD

- ▶ Eigen-Value Decomposition (EVD)
- ▶ Singular Value Decomposition (SVD)

Singular Value Decomposition

SVD

An $n \times p$ matrix \mathbf{X} can be decomposed as:

$$\mathbf{X} = \mathbf{U}\mathbf{\Lambda}\mathbf{V}'$$

where

- ▶ \mathbf{U} is a $n \times p$ column **orthonormal** matrix containing the left singular vectors
- ▶ $\mathbf{\Lambda}$ is a $p \times p$ **diagonal** matrix containing the singular values of \mathbf{X}
- ▶ \mathbf{V} is a $p \times p$ column **orthonormal** matrix containing the right singular vectors

svd() in R

svd() function

R provides the function `svd()` to perform a singular value decomposition of a given matrix

svd() output

A list with the following components

- `d` a vector containing the singular values
- `u` a matrix whose columns contain the left singular vectors
- `v` a vector whose columns contain the right singular vectors

SVD example in R

```
# X matrix
set.seed(22)
X = matrix(rnorm(20), 5, 4)

# singular value decomposition
SVD = svd(X)

# elements returned by svd()
names(SVD)

## [1] "d" "u" "v"

# vector of singular values
(d = SVD$d)

## [1] 3.9516 2.0224 1.4748 0.4324
```

```
# matrix of left singular vectors
(U = SVD$u)

##           [,1]      [,2]      [,3]      [,4]
## [1,] -0.4251 -0.53913 -0.7233  0.009794
## [2,]  0.5269 -0.76863  0.2860  0.056100
## [3,]  0.5753  0.05000 -0.4421  0.131072
## [4,]  0.2215  0.05273 -0.1702 -0.951234
## [5,] -0.4021 -0.33655  0.4131 -0.273371

# matrix of right singular vectors
(V = SVD$v)

##           [,1]      [,2]      [,3]      [,4]
## [1,]  0.5708 -0.7407  0.33863  0.1043
## [2,] -0.2742 -0.5295 -0.76797  0.2338
## [3,]  0.2772  0.3206 -0.04462  0.9046
## [4,]  0.7226  0.2612 -0.54181 -0.3408
```

SVD example in R (con't)

```
# U orthonormal (U'U = I)
```

```
t(U) %*% U
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,]  1.000e+00  1.388e-16  2.776e-17  0.000e+00
## [2,]  1.388e-16  1.000e+00 -2.776e-17 -8.327e-17
## [3,]  2.776e-17 -2.776e-17  1.000e+00  5.551e-17
## [4,]  0.000e+00 -8.327e-17  5.551e-17  1.000e+00
```

```
# V orthonormal (V'V = I)
```

```
t(V) %*% V
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,]  1.000e+00 -1.110e-16 -5.551e-17  1.110e-16
## [2,] -1.110e-16  1.000e+00  8.327e-17  1.943e-16
## [3,] -5.551e-17  8.327e-17  1.000e+00 -8.327e-17
## [4,]  1.110e-16  1.943e-16 -8.327e-17  1.000e+00
```

```
# X equals U D V'
```

```
U %*% diag(d) %*% t(V)
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,] -0.5121  1.85809 -0.76391 -0.9222
## [2,]  2.4852 -0.06603  0.08196  0.8616
## [3,]  1.0078 -0.16276  0.74303  2.0029
## [4,]  0.2928 -0.19986 -0.08402  0.9366
## [5,] -0.2090  0.30056 -0.79289 -1.6157
```

```
# compare to X
```

```
X
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,] -0.5121  1.85809 -0.76391 -0.9222
## [2,]  2.4852 -0.06603  0.08196  0.8616
## [3,]  1.0078 -0.16276  0.74303  2.0029
## [4,]  0.2928 -0.19986 -0.08402  0.9366
## [5,] -0.2090  0.30056 -0.79289 -1.6157
```